

OwnRAG parsing pipeline

A PDF is opened page by page so every citation can carry a page number. Text is extracted per page and han

Chunking

The naive method groups paragraphs up to 512 tokens. The laws method splits on article boundaries instead, w

Indexing

Each chunk is tokenised, embedded, and written to both the inverted index and the vector index. A chunk is seen

OwnRAG parsing pipeline

A PDF is opened page by page so every citation can carry a page number. Text is extracted per page and han

Chunking

The naive method groups paragraphs up to 512 tokens. The laws method splits on article boundaries instead, w

Indexing

Each chunk is tokenised, embedded, and written to both the inverted index and the vector index. A chunk is seen

OwnRAG parsing pipeline

A PDF is opened page by page so every citation can carry a page number. Text is extracted per page and han

Chunking

The naive method groups paragraphs up to 512 tokens. The laws method splits on article boundaries instead, w

Indexing

Each chunk is tokenised, embedded, and written to both the inverted index and the vector index. A chunk is seen

OwnRAG parsing pipeline

A PDF is opened page by page so every citation can carry a page number. Text is extracted per page and han

Chunking

The naive method groups paragraphs up to 512 tokens. The laws method splits on article boundaries instead, w

Indexing

Each chunk is tokenised, embedded, and written to both the inverted index and the vector index. A chunk is seen

OwnRAG parsing pipeline

A PDF is opened page by page so every citation can carry a page number. Text is extracted per page and han

Chunking

The naive method groups paragraphs up to 512 tokens. The laws method splits on article boundaries instead, w

Indexing

Each chunk is tokenised, embedded, and written to both the inverted index and the vector index. A chunk is seen

OwnRAG parsing pipeline

A PDF is opened page by page so every citation can carry a page number. Text is extracted per page and han

Chunking

The naive method groups paragraphs up to 512 tokens. The laws method splits on article boundaries instead, w

Indexing

Each chunk is tokenised, embedded, and written to both the inverted index and the vector index. A chunk is seen

OwnRAG parsing pipeline

A PDF is opened page by page so every citation can carry a page number. Text is extracted per page and han

Chunking

The naive method groups paragraphs up to 512 tokens. The laws method splits on article boundaries instead, w

Indexing

Each chunk is tokenised, embedded, and written to both the inverted index and the vector index. A chunk is seen

OwnRAG parsing pipeline

A PDF is opened page by page so every citation can carry a page number. Text is extracted per page and han

Chunking

The naive method groups paragraphs up to 512 tokens. The laws method splits on article boundaries instead, w

Indexing

Each chunk is tokenised, embedded, and written to both the inverted index and the vector index. A chunk is seen

OwnRAG parsing pipeline

A PDF is opened page by page so every citation can carry a page number. Text is extracted per page and han

Chunking

The naive method groups paragraphs up to 512 tokens. The laws method splits on article boundaries instead, w

Indexing

Each chunk is tokenised, embedded, and written to both the inverted index and the vector index. A chunk is seen

OwnRAG parsing pipeline

A PDF is opened page by page so every citation can carry a page number. Text is extracted per page and han

Chunking

The naive method groups paragraphs up to 512 tokens. The laws method splits on article boundaries instead, w

Indexing

Each chunk is tokenised, embedded, and written to both the inverted index and the vector index. A chunk is seen

OwnRAG parsing pipeline

A PDF is opened page by page so every citation can carry a page number. Text is extracted per page and han

Chunking

The naive method groups paragraphs up to 512 tokens. The laws method splits on article boundaries instead, w

Indexing

Each chunk is tokenised, embedded, and written to both the inverted index and the vector index. A chunk is seen

OwnRAG parsing pipeline

A PDF is opened page by page so every citation can carry a page number. Text is extracted per page and han

Chunking

The naive method groups paragraphs up to 512 tokens. The laws method splits on article boundaries instead, w

Indexing

Each chunk is tokenised, embedded, and written to both the inverted index and the vector index. A chunk is seen